

RISC-V Lab

Ex4: On Chip Interconnects

Content

1. Basics
2. Pipelining
3. Decouple
4. Topologies
5. Multi Master
6. Implementation
7. Examples

Basics

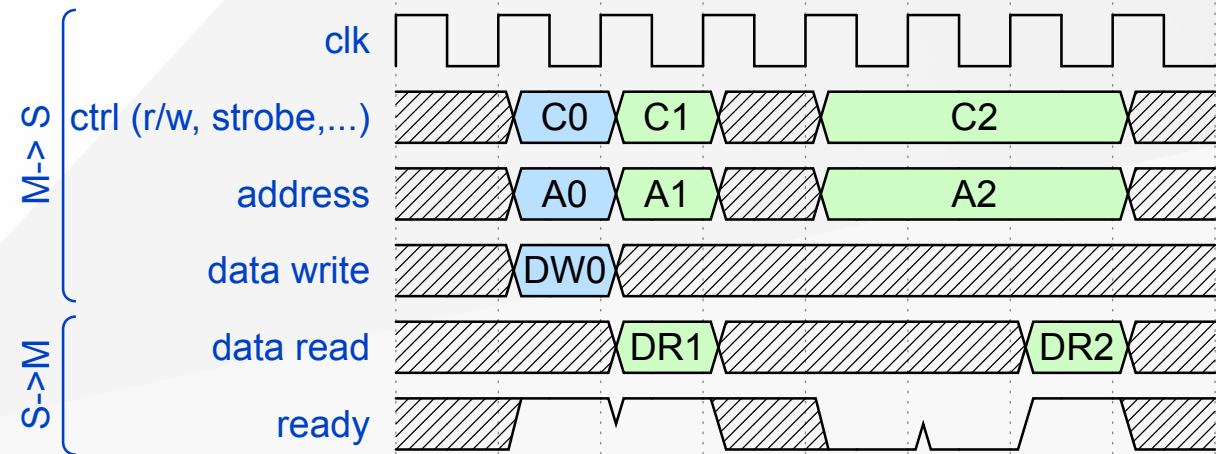
- requirement: *exchange data* between N on chip components.
 - issues: on chip delays, $O(N^2)$ for full orthogonality
 - performance metrics: throughput & latency
- **common logical model** of all components: hardware resources (registers, memory) mapped into "flat" address space
- primarily standardized **interface**, not topology
- maximizes **reuse** ("lego") of
 - know how
 - modules
 - test, debugging & profiling infrastructure

Basics: Nodes

Classes of nodes

- master / host: initiates transfer
- slave / device: determines speed of transfer

Slave can stop a transfer anytime (ready, ack ...)
=> master / OCI blocked
=> loss of throughput (& latency)



Prevent blockade of master & OCI

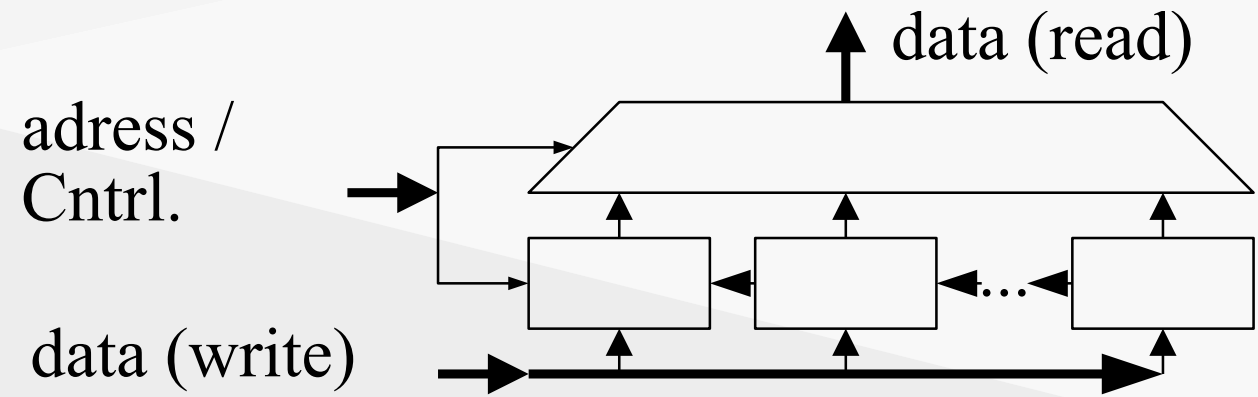
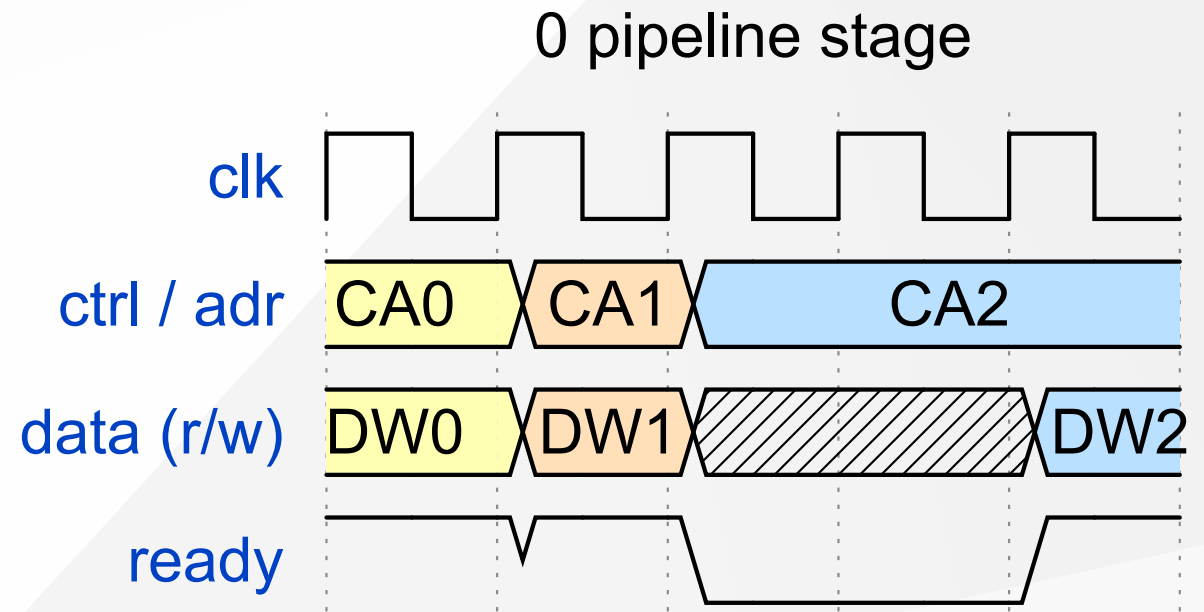
docs/exercises/res/ex4_tlul_a.svg

1. Master provides additional information to the slave
e.g. amount & addresses next ("burst" access)
2. Split transfer
Master releases bus for new transfers during the wait period and the transfer is resumed later (often not feasible / too complex)
3. Pipelining
4. Decouple request from action (read/write)
-> message based (NoC)

Pipelining

0 pipeline stages

- 0..1 transfers in flight
- latency: 0 cycles
- combinational loop
master -> slave -> master
(adr/ctrl -> "ready" / data(r))



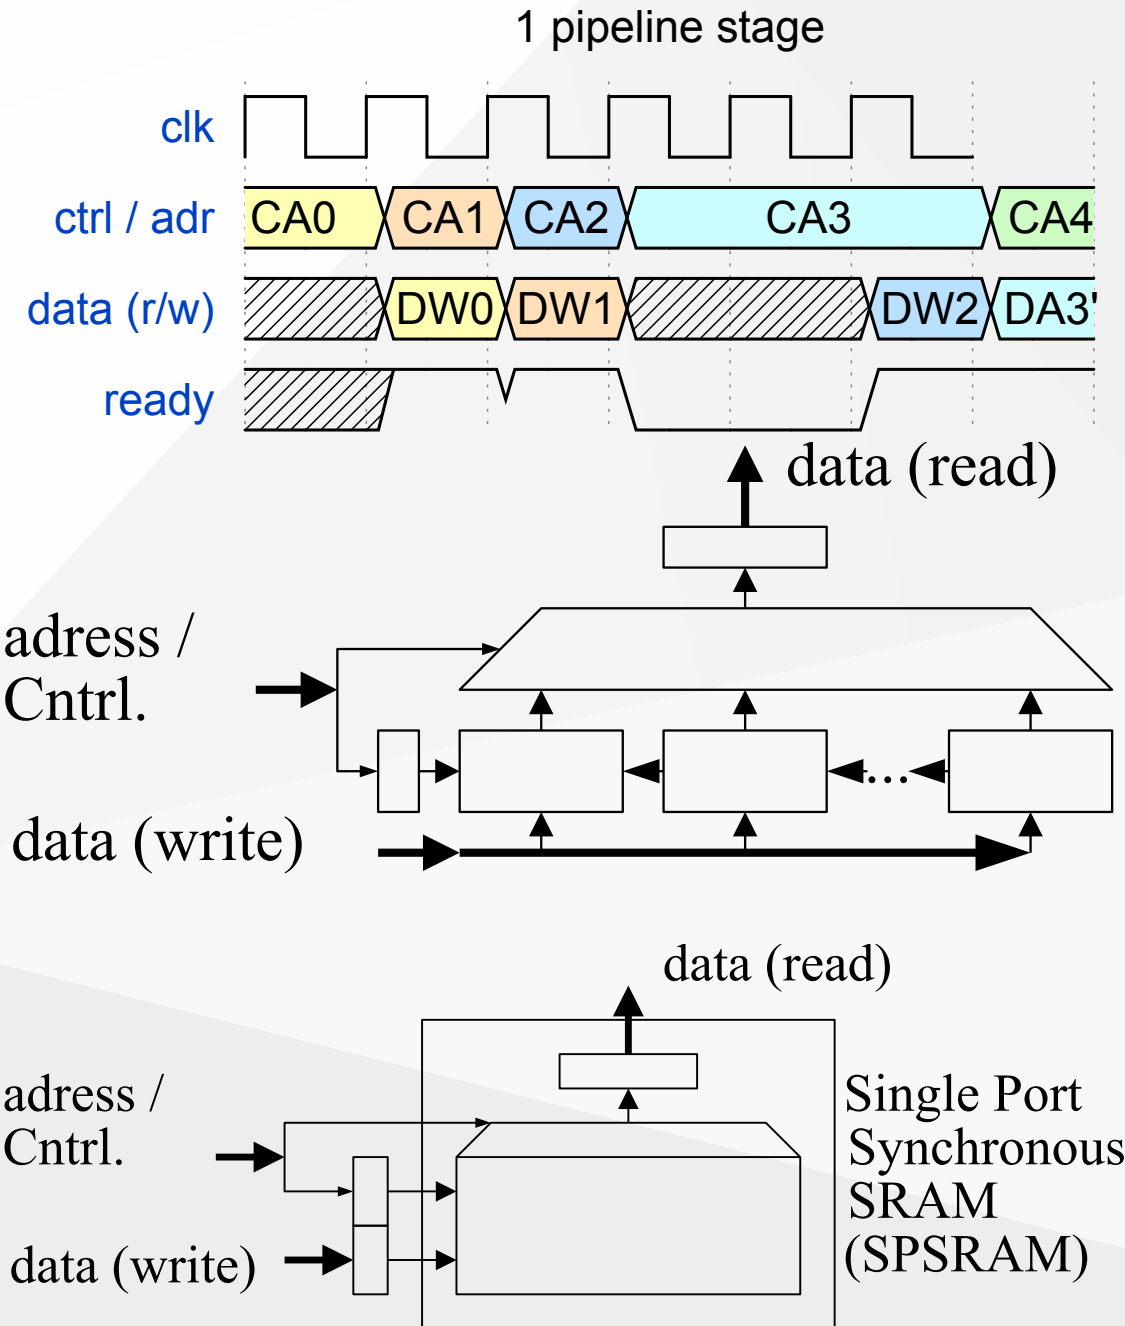
Pipelining

1 pipeline stage

- 0..2 transfers in flight
- latency: 0 cycle
- hold by slaves holds the *next* transfer

1 read / 0 write pipeline stages

- used by SPSRAM

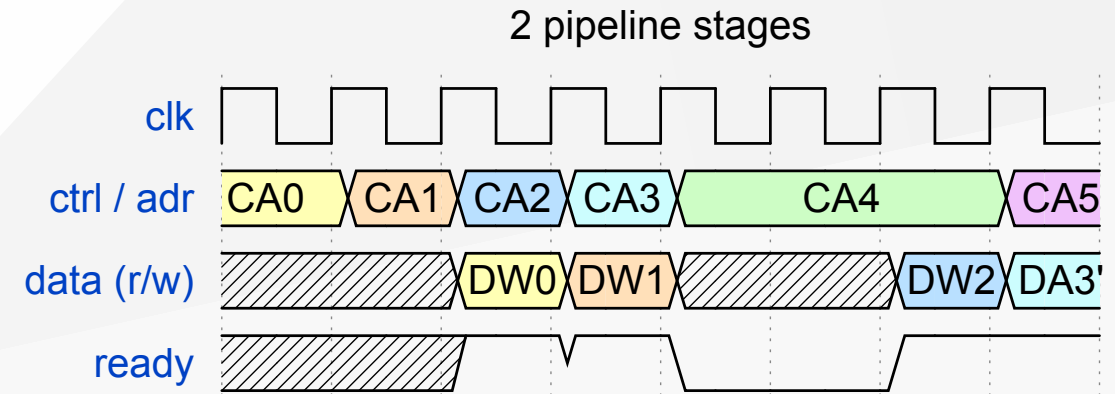


Pipelining: Conversion 1r/0w -> 1r/1w

Pipelining: 2 stages

- very complex (esp. crossbars)
- for some cases
 - too long (e.g. CPU 2 SRAM)
 - too short (e.g. CPU 2 far away peripheral)

=> need "variable" pipeline depth



Decouple

decouple action (read/write) from bus transaction

- "bus read/write" => request message + response message
- independent channels for request & response
- handshaking for messages, not actions

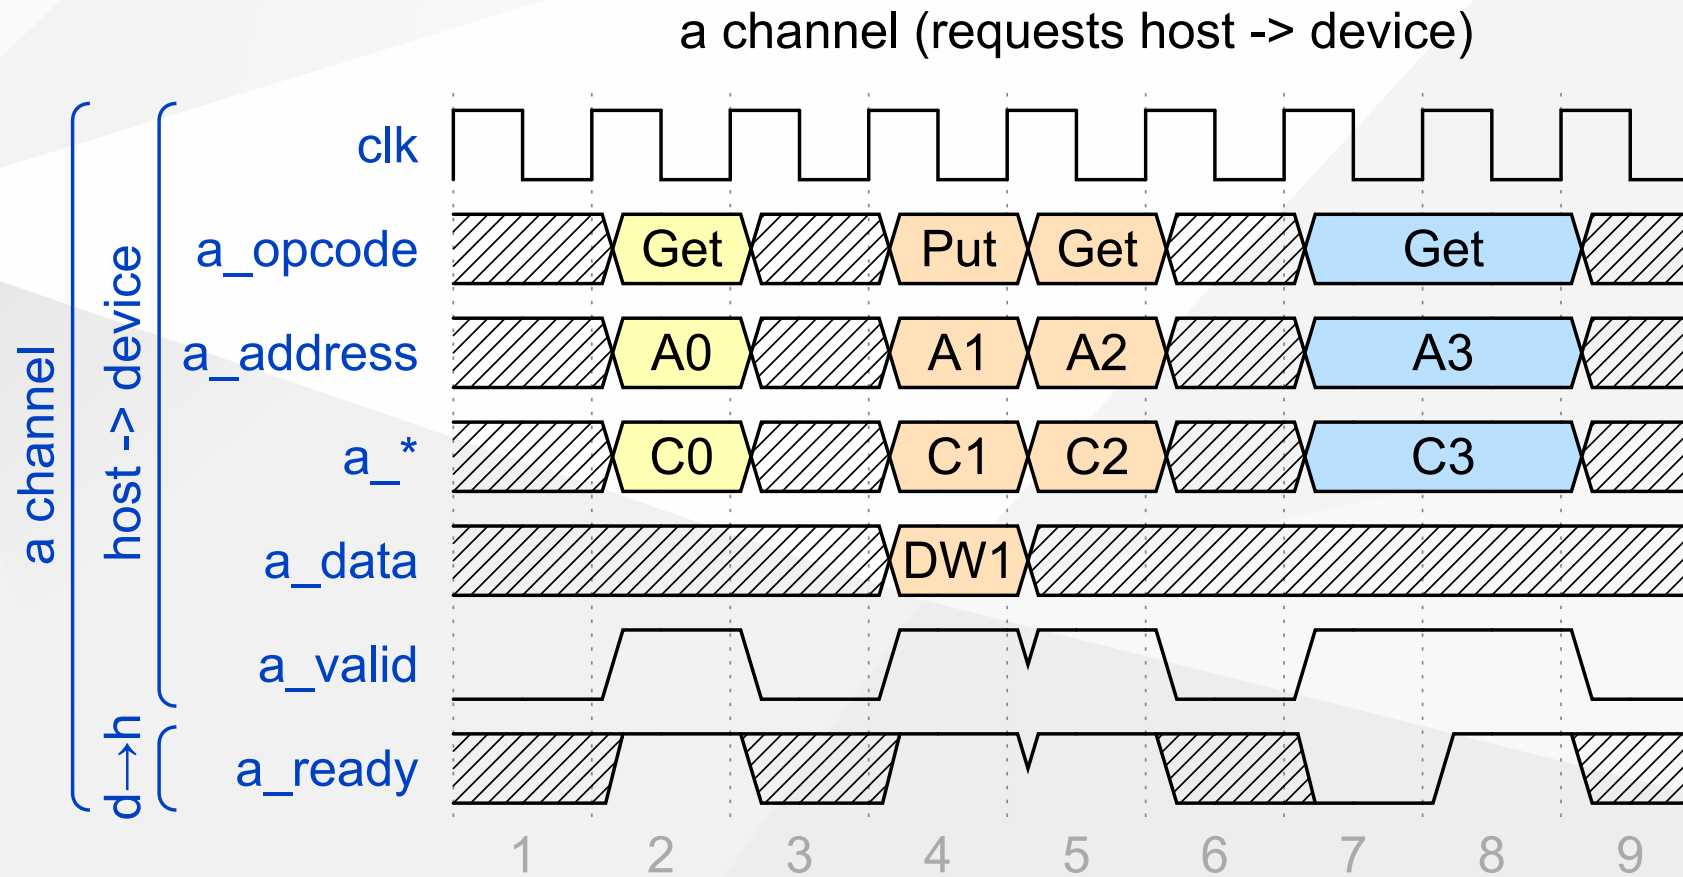
=> bus latencies do not impact throughput

=> $0..∞$ simultaneous transactions

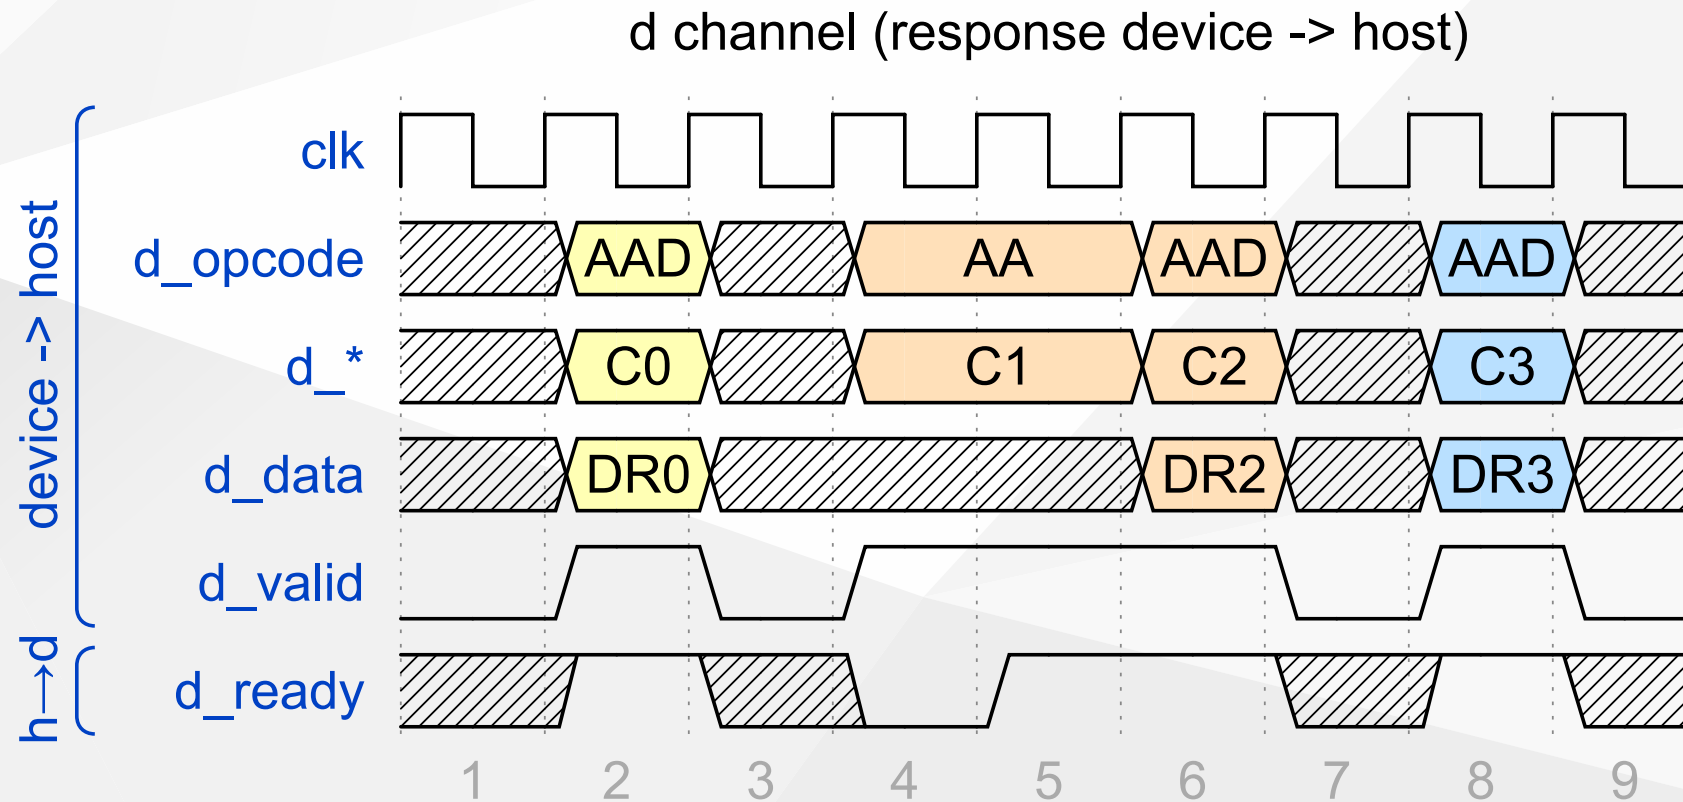
Examples

- ARM AMBA **AXI** (Advanced eXtensible Interface)
- Tile Link

Decouple: TL-UL: a channel



Decouple: TL-UL: d channel



Decouple: TL-UL: Rules


- goal: prevent deadlock & combinational loop
- "4.1 Flow Control Rules"
 - valid must never depend on ready
(no comb. path from ready to any control / data signal)
 - ...
- "4.2 Deadlock Freedom"
 - $\text{prio}(\text{response}) > \text{prio}(\text{request})$
 - ...

Please read the TL-UL specification!

Topologies: Extremes

docs/exercises/res/ex4_tlul_a.svg

Topology	connections	simultaneous	comment
(shared) bus	M:S, 1:S	1	e.g. (legacy) PCB
crossbar / switch	M:S	$\min(M,S)$	completely connected
point 2 point	1:1	1	



Topologies: Hierarchical

- many intermediates topologies: ring, cube , ...
- hierarchical structures using (sparse) crossbars
 - sparse = subset of (M,S)
 - extremes: 1:N ("bus"), M:1

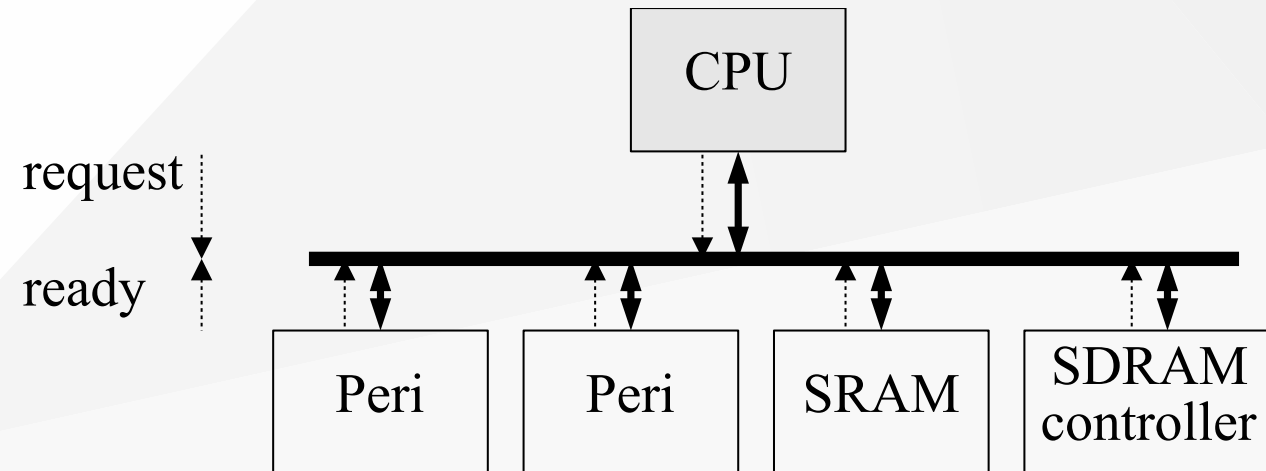
The diagram illustrates the internal structure of the **rvlab_core**. It features a central **rvlab_cpu** block with **instr. fetch** and **load/store** units. This CPU is connected to an **xbar_main** (crossbar) and an **xbar_peri** (peripheral crossbar). The **xbar_main** connects the CPU to **ram_main** and **rvlab_debug**. The **xbar_peri** connects the CPU to **rv_timer** and **regdemo**. A **student** block is connected to the **rvlab_core** via **user_io_i** and **user_io_o** signals. The **student** also interacts with the **rvlab_core** through **tl_device_fast_i** and **tl_device_fast_o** signals. The **rvlab_core** is connected to an external **user_io** block via **user_io_i** and **user_io_o** signals. The **rvlab_core** also has a **tl_ddd_o** and **tl_ddd_i** signal, and a **tl_ddd_ctrl_o** and **tl_ddd_ctrl_i** signal. The **rvlab_core** is connected to an external **user_io** block via **user_io_i** and **user_io_o** signals. The **rvlab_core** also has a **tl_ddd_o** and **tl_ddd_i** signal, and a **tl_ddd_ctrl_o** and **tl_ddd_ctrl_i** signal. The **rvlab_core** is connected to an external **user_io** block via **user_io_i** and **user_io_o** signals. The **rvlab_core** also has a **tl_ddd_o** and **tl_ddd_i** signal, and a **tl_ddd_ctrl_o** and **tl_ddd_ctrl_i** signal.



Single Master

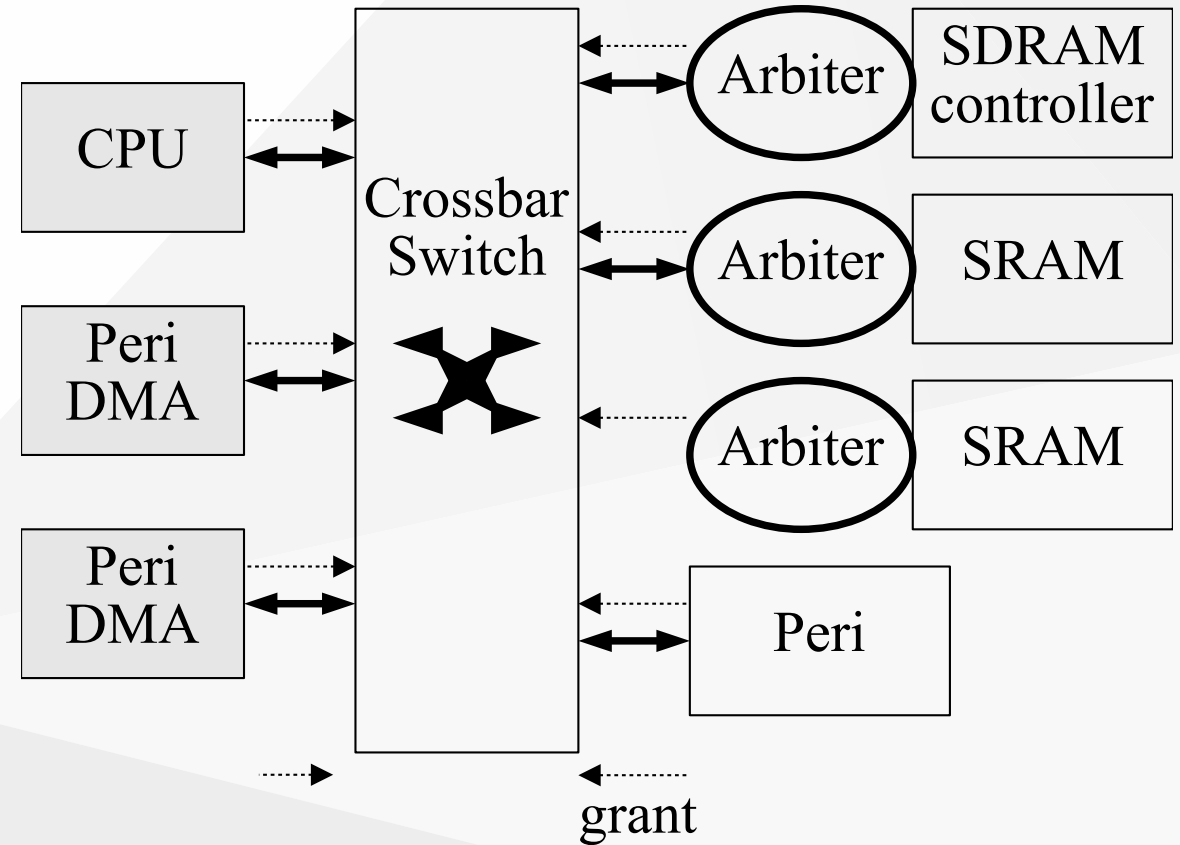
1. PERI -> CPU
 2. /3. (CPU <-> RAM)
 3. CPU -> PERI
- **every** word moved by the CPU at least twice
 - interconnect latency to peri directly slows down CPU

Ex.: SPI without DMA handled
by 400 MHz R7 ARM in an
SoC => <20 Mbit



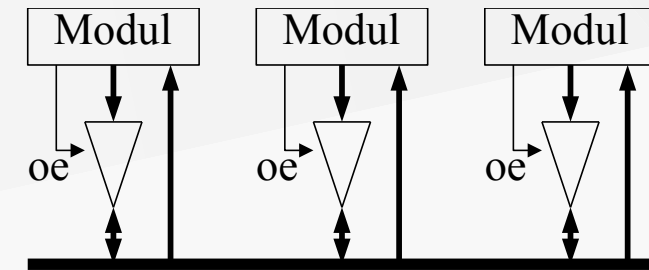
Multi Master

1. peri -> RAM
2. CPU switches pointers
3. RAM -> peri
 - peri works on (complex) data structures (e.g. descriptors, linked lists, instructions)
 - CPU works in parallel
 - interconnect latency to peri hidden from CPU



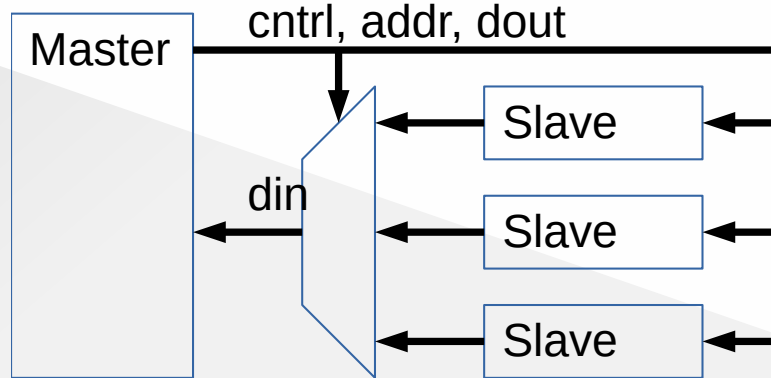
Implementation: Tristate

- lines (data, address, ...) used bidirectionally
- used in: PCB design
 - orders of magnitude fewer components
 - limits on package pin count and routing
 - mostly for medium (SQL, PSRAM) to slow (I2C) components
- on chip: => **Tristate not suitable.**
 - limitations order of magnitude higher
 - high design effort, (e.g. 2 phase clk, asyn. delays, bus holders ...)
 - not supported by EDA tools

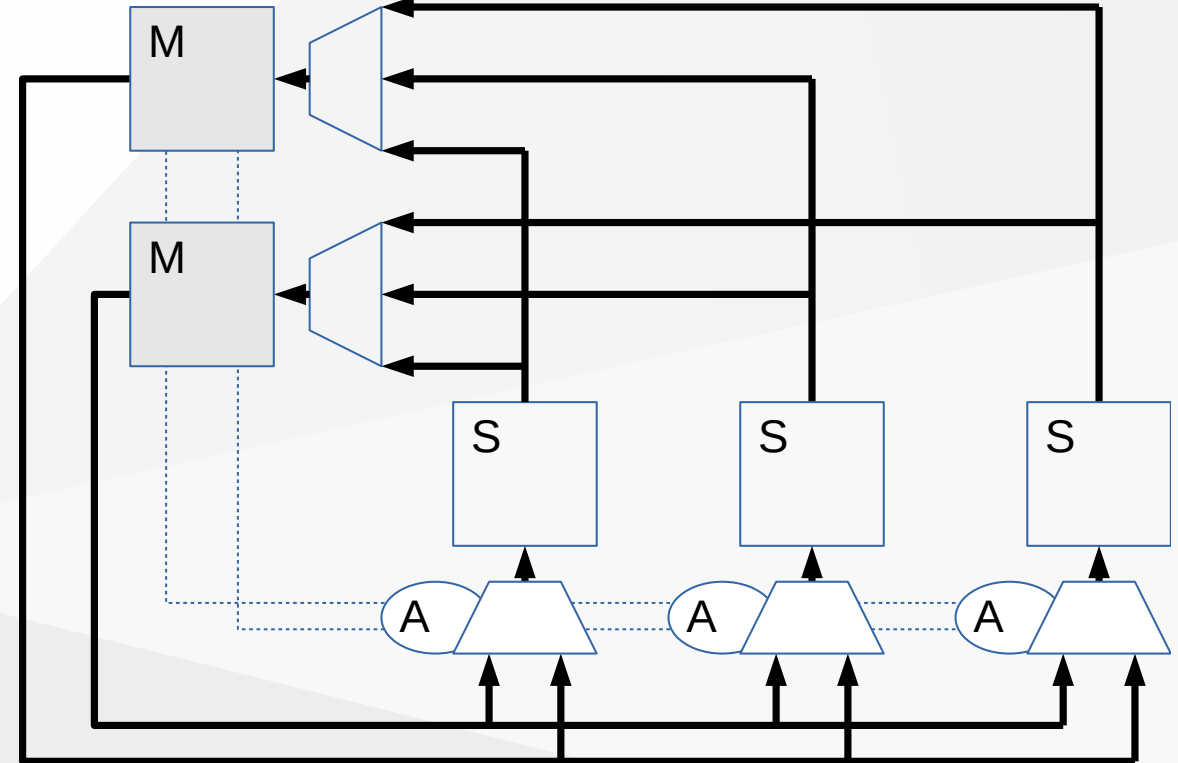


Implementation: Multiplexed

**Multiplexed implementation
“bus” (1:S, “single master”)**

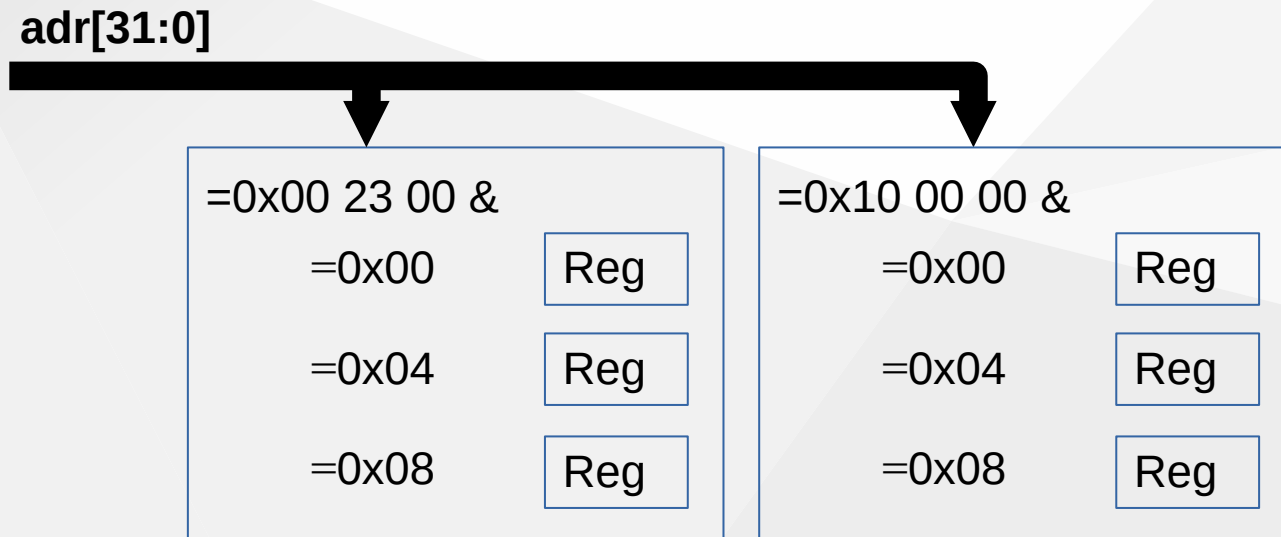


**Multiplexed implementation
full crossbar (M:S, “multi master”)**



Implementation: Flat address decoding

- every module places its registers at absolute addresses in total address space
- (unnecessary) high delay & gate count
- non existent orthogonality



Impl.: Hierarchical address decoding

- every module places its registers relative to address 0, i.e. it only decodes the **lowermost** address bits **used by itself**
- a further signal indicates a module access

